

Cadence Catalog Intelligence

SEO & AI Search Launch Plan

What we shipped, what to expect, and the step-by-step playbook for getting cadence-ci.com indexed by Google, Bing, ChatGPT, and Perplexity.

Prepared for the Cadence team · May 2026

1. What we just shipped

We rolled out the technical foundation that lets search engines and AI assistants (Google, Bing, ChatGPT, Perplexity, Claude) actually find, understand, and recommend cadence-ci.com when people search for music catalog software, publisher analytics, or rights management tools.

Plain-English summary

- **Title & description tags** on every public page so search results show a clean, branded preview that says *Cadence Catalog Intelligence* (not just “Cadence” — important because Cadence Design Systems is a different, much larger company).
- **Open Graph & Twitter cards** so when anyone shares a Cadence link on LinkedIn, X, iMessage, or Slack, it expands into a polished card with our logo, tagline, and image.
- **A purpose-built share image** (1200×630 PNG) that appears as the preview thumbnail.
- **Structured data (JSON-LD)** describing Cadence as an Organization and a SoftwareApplication with three pricing tiers (Basic \$29.99, Pro \$99.99, Enterprise \$499.99). This is what AI assistants read when someone asks ChatGPT or Perplexity “what does Cadence Catalog Intelligence do?”
- **sitemap.xml** listing every public URL so Google can crawl them efficiently.
- **robots.txt** telling crawlers which areas to index (public pages) and which to ignore (login pages, dashboards, internal admin).
- **A hidden long-form page** at /what-is-cadence with the full keyword-rich product description — visible to Google & AI crawlers but **not linked anywhere** on the visible site, so competitors don’t see our roadmap when they scroll the homepage.

Why we hid the long-form page

Originally we put the full feature breakdown (catalog management, valuation, royalty processing, rights administration, sync placements, creator portals, streaming intelligence) right on the homepage so search engines could crawl it. The team flagged that this also exposed our entire product strategy to anyone who scrolled — competitors included. The fix: move that content to a dedicated page (/what-is-cadence) that’s listed in the sitemap (so Google, ChatGPT, Perplexity, etc. find and index it) but isn’t linked from the homepage, footer, or nav. The walk-in visitor sees the vague positioning we want; the AI search engines get the full picture.

TECHNICAL — Implementation files

frontend/index.html — base meta tags + JSON-LD blocks (Organization, SoftwareApplication with 3 offers).
frontend/src/components/SEO.jsx — react-helmet-async per-route title/description/canonical override component.
frontend/src/pages/WhatIsCadencePage.jsx — hidden long-form page (unlinked, sitemap-only).
frontend/src/App.jsx — route definition for /what-is-cadence.
frontend/public/sitemap.xml, robots.txt, og-image.png — static SEO assets.

2. What to expect — the realistic timeline

SEO is not instant. Here's what actually happens after we publish:

What	When	Notes
Pages live & crawlable	Immediately	As soon as the deploy finishes (~1–2 min).
Google discovers /what-is-cadence	1–7 days	Faster if you submit the sitemap manually (Step 3 below).
Google indexes & starts ranking	1–4 weeks	Use Search Console “Request Indexing” to fast-track to 24–48 hours.
Bing / DuckDuckGo indexing	3–14 days	Bing Webmaster Tools speeds this up.
Perplexity surfacing the site	1–2 weeks	Perplexity uses Bing’s index, so once Bing has it, Perplexity does too.
GPTBot / ChatGPT picking it up	2–8 weeks	OpenAI’s crawler has its own slower cadence; nothing you can do to speed it up.
Meaningful organic traffic	2–3 months	Once Google has crawled, indexed, and started ranking pages.

Bottom line: Plan for ~2 weeks before you start seeing search results, ~6 weeks before AI assistants reliably mention us, and ~3 months before SEO becomes a meaningful traffic channel. The work below is what cuts those timelines in half.

3. Google Search Console — step by step

Search Console is Google's free tool for telling them "this site exists, please crawl it." It also shows you exactly which search queries are bringing people to the site once you start ranking. Total setup time: ~10 minutes.

3.1 — Set up the property (one-time, ~5 min)

- 1 Go to **search.google.com/search-console** and sign in with a Google account. Use a business account if possible — whoever owns this account owns the SEO data.
- 2 Click **Add Property** → choose **Domain** (preferred — it covers cadence-ci.com plus all subdomains and both http/https).
- 3 Google will give you a **TXT record** that looks like `google-site-verification=AbCdEf1234...`
- 4 Add that TXT record in your DNS provider's control panel (Cloudflare, Namecheap, GoDaddy, Replit Domains, wherever cadence-ci.com is registered). The exact path is usually *DNS* → *Add Record* → *Type: TXT* → *Name: @ (or blank)* → *Value: paste the string Google gave you*.
- 5 Wait 5–15 minutes for DNS to propagate, then click **Verify** in Search Console.

If you can't do the Domain verification (some registrars make it tricky), the easier fallback is the **URL prefix** option using HTML file upload — but Domain is better long-term because it captures every variant of the URL.

3.2 — Submit the sitemap (one-time, 30 sec)

- 1 In Search Console's left sidebar, click **Sitemaps**.
- 2 In the "Add a new sitemap" box, type just: `sitemap.xml`
- 3 Click **Submit**. Done.

Google will now check that file regularly and discover all 11 URLs we listed there — including the hidden /what-is-cadence page.

3.3 — Request indexing for each important URL

This is the highest-leverage step. Without it, Google may take weeks to crawl a new page; **with it, usually within 24–48 hours**.

- 1 At the very top of Search Console there's a search bar that says "Inspect any URL in cadence-ci.com".
- 2 Paste the full URL, e.g. `https://cadence-ci.com/what-is-cadence`
- 3 Press Enter. Google fetches the URL live (takes ~30 sec).
- 4 If it says "URL is not on Google" (expected for a brand-new page), click **Request Indexing**.
- 5 Google queues a priority crawl. Repeat for each URL on the checklist below.

There's a daily quota of ~10–12 indexing requests, so spread the list below across two days if you hit the limit.

URL checklist — submit these in this priority order

- <https://cadence-ci.com/> — the homepage
- <https://cadence-ci.com/what-is-cadence> — the hidden SEO page (most important)
- <https://cadence-ci.com/about>
- <https://cadence-ci.com/help>
- <https://cadence-ci.com/careers>
- <https://cadence-ci.com/investors>

4. Bing, DuckDuckGo, Perplexity & ChatGPT

4.1 — Bing Webmaster Tools (10 min, optional but worth it)

Bing has its own crawler — and Bing's index also feeds DuckDuckGo, Yahoo, Ecosia, and Perplexity. So getting into Bing is essentially “getting into every non-Google search experience.”

- 1 Go to **bing.com/webmasters** and sign in (Microsoft account).
- 2 Click **Import from Google Search Console** — this is the magic button. It copies your verified property and sitemap directly from Google so you don't have to re-do DNS verification.
- 3 Once imported, use Bing's **URL Submission** tool the same way as Google's Request Indexing. Bing's daily quota is much higher (~10,000) so submit everything at once.

4.2 — AI assistants (ChatGPT, Claude, Perplexity)

There's no “submit my site to ChatGPT” button. AI assistants surface Cadence in two ways:

- **Live web search** (Perplexity always, ChatGPT & Claude when web access is on). These tools query Google or Bing in real time, so once we're indexed there, we're reachable here. Perplexity in particular leans heavily on Bing.
- **Trained-in knowledge** (the model “knows” about Cadence without searching). This requires the model's training data to include our pages. OpenAI's GPTBot, Anthropic's ClaudeBot, and Google's extended crawlers all visit indexed sites on their own schedule (typically every 4–12 weeks). Our robots.txt does not block any of these crawlers, so this happens automatically over time.

The JSON-LD structured data we shipped is specifically designed for AI consumption — it says “Cadence Catalog Intelligence is a SoftwareApplication, here's what category, here's the pricing” in machine-readable form. This is the single biggest thing you can do to help AI assistants describe us correctly.

TECHNICAL — How AI crawlers see the site

Both static meta tags (rendered server-side in `frontend/index.html`) and per-route tags (rendered client-side via `react-helmet-async`) are visible to modern crawlers. Google, Bing, GPTBot, and ClaudeBot all execute JavaScript before extracting content, so the SPA architecture is not a blocker. The `SoftwareApplication` JSON-LD with three `Offer` objects (Basic/Pro/Enterprise pricing) is the canonical schema for AI assistants asked “how much does X cost?” Keep it accurate — if pricing changes, update `frontend/index.html`.

5. How to verify it's working

5.1 — Right now (before any indexing)

- 1 Visit <https://cadence-ci.com/what-is-cadence> directly — confirm the rich descriptive copy renders.
- 2 Open <https://cadence-ci.com/>, scroll the entire homepage, click around the footer. Confirm **nothing links to /what-is-cadence**. (That's the whole point — it's for crawlers, not visitors.)
- 3 Visit <https://cadence-ci.com/sitemap.xml> — confirm /what-is-cadence is listed.
- 4 Visit <https://cadence-ci.com/robots.txt> — confirm it's not blocked.
- 5 Paste the URL into Google's free **Rich Results Test** (search.google.com/test/rich-results) — should show the Organization and SoftwareApplication structured data parsing correctly.

5.2 — Weekly, in Search Console

- **Pages report** — how many of our URLs Google has indexed. Should grow from 1 → 6 → 11+ over 2–3 weeks.
- **Performance tab** — once we have impressions, this shows the actual queries people typed to find us, our average position, and click-through rate.
- **Coverage / Errors** — flags any pages Google couldn't crawl. Address these promptly.

5.3 — The smell test (3–4 weeks after publishing)

Open Google in an incognito window and try these searches. If we're showing up, SEO is working:

- *cadence catalog intelligence* — should be result #1 within a week.
- *cadence-ci.com* — should return our pages.
- *music catalog management software*
- *publisher analytics platform*
- *music catalog valuation tool*
- *rights management software for music publishers*

Then ask the same questions to ChatGPT, Claude, and Perplexity. Within 4–6 weeks they should be able to describe Cadence Catalog Intelligence (what it does, who it's for, the pricing tiers) without us being explicitly mentioned in the prompt.

6. Owner checklist

Suggested division of labor — assign one owner per item.

This week

- ☐ Set up Google Search Console (Section 3.1)
- ☐ Submit sitemap.xml in Search Console (Section 3.2)
- ☐ Request indexing for all 6 URLs in the checklist (Section 3.3)
- ☐ Set up Bing Webmaster Tools via the Import-from-Google button (Section 4.1)
- ☐ Run the 5 verification checks in Section 5.1

Within 30 days

- ☐ Check Search Console weekly — track Pages indexed and impressions
- ☐ Run the “smell test” searches in Section 5.3
- ☐ Ask ChatGPT & Perplexity “what is Cadence Catalog Intelligence?” — track answer quality
- ☐ If a key URL still isn’t indexed after 14 days, re-request indexing

Quarterly

- ☐ Review which queries are bringing the most impressions/clicks (Search Console → Performance)
- ☐ Update sitemap.xml lastmod dates if any pages have been substantively rewritten
- ☐ Refresh JSON-LD pricing in frontend/index.html if tier prices have changed
- ☐ Consider adding 1–2 new content pages (use cases, customer stories) to the sitemap as the product matures

Appendix — Technical reference

For developers maintaining the SEO setup. Skip this section if you're not working with the codebase.

File map

File	Purpose
frontend/index.html	Base meta tags, OG/Twitter cards, JSON-LD (Organization + SoftwareApplication w/ 3 offers). Loaded for every route.
frontend/src/components/SEO.jsx	react-helmet-async wrapper. Per-route override of <title>, meta description, canonical URL, OG/Twitter tags.
frontend/src/main.jsx	Mounts <HelmetProvider> so per-route SEO components work.
frontend/src/pages/WhatIsCadencePage.jsx	Hidden long-form SEO page. Unlinked. Discoverable only via direct URL or sitemap.
frontend/src/App.jsx	Route definition for /what-is-cadence in the unauthenticated routes block.
frontend/public/sitemap.xml	Static sitemap listing 11 public URLs. Update lastmod when pages change.
frontend/public/robots.txt	Allow public, disallow authenticated app surfaces. References sitemap.
frontend/public/og-image.png	1200x630 share preview image. Referenced from index.html OG/Twitter blocks.

Key architectural decisions

- **Brand-leading titles.** Every <title> begins with “Cadence Catalog Intelligence” to disambiguate from Cadence Design Systems (a much larger company) in search results.
- **Hardcoded fallback in index.html.** Static title/description/JSON-LD are present in the served HTML so crawlers without JS execution still see canonical metadata. Per-route react-helmet-async tags override at runtime.
- **Hidden page architecture for /what-is-cadence.** The page is publicly accessible (not gated, not cloaked, not user-agent-sniffed — that would violate Google's quality guidelines). It is simply unlinked from the visible UI. Discoverability comes from sitemap.xml inclusion. This satisfies SEO/AI discoverability without exposing the full feature breakdown to walk-in visitors.

- **JSON-LD pricing.** Three Offer objects (Basic 29.99, Pro 99.99, Enterprise 499.99 USD) in the SoftwareApplication block. Update these in frontend/index.html when pricing changes; this is the canonical answer AI assistants give when asked about cost.
- **Disallow authenticated routes.** robots.txt blocks /dashboard, /catalog, /roster, /contracts, /admin, /internal, etc. so private app surfaces never appear in search results.

Maintenance triggers

- **Pricing change** → update Offer prices in JSON-LD inside frontend/index.html.
- **New public page** → add <Route> in App.jsx, mount <SEO> component on the page, add <url> entry to sitemap.xml.
- **Brand or tagline change** → update the static <title>, meta description, OG tags in frontend/index.html, plus the default fullTitle in SEO.jsx.
- **Pricing tier added/removed** → sync both the JSON-LD offers and any pricing page copy (currently no pricing page is built — flagged for a future task).

Questions? See replit.md (“SEO & AI Search Visibility” section) for the current architectural summary, or check the commit history on Tasks #202 and #203 for context on why each piece was built the way it was.